

Online learning, minimizing regret, and combining expert advice.

Maria Florina (Nina) Balcan

04/29/2019

- “The weighted majority algorithm” N. Littlestone & M. Warmuth
- “Online Algorithms in Machine Learning” (survey) A. Blum

Motivation


Many situations involve repeated decision making in an uncertain environment.

- Deciding how to invest your money (buy or sell stocks)
- What route to drive to work each day
- Playing repeatedly a game against an opponent with unknown strategy

We will study:

Learning algos for such settings with connections to game theoretic notions of equilibria

Online learning, minimizing regret, and combining expert advice.



Using "expert" advice

Assume we want to predict the stock market.

- We solicit n "experts" for their advice.
 - Will the market go up or down?
- We then want to use their advice somehow to make our prediction. E.g.,

Expt 1	Expt 2	Expt 3	neighbor's dog	truth
down	up	up	up	up
down	up	up	down	down
...

Can we do nearly as well as best in hindsight?



Note: "expert" \equiv someone with an opinion.

[Not necessarily someone who knows anything.]



Formal model

- There are n experts.
- For each round $t=1, 2, \dots, T$
 - Each expert makes a prediction in $\{0, 1\}$
 - The learner (using experts' predictions) makes a prediction in $\{0, 1\}$
 - The learner observes the actual outcome. There is a mistake if the predicted outcome is different from the actual outcome.
The learner gets to update his hypothesis.

Can we do nearly as well as best in hindsight?



Formal model

- There are n experts.
- For each round $t=1, 2, \dots, T$
 - Each expert makes a prediction in $\{0,1\}$
 - The learner (using experts' predictions) makes a prediction in $\{0,1\}$
 - The learner observes the actual outcome. There is a mistake if the predicted outcome is different from the actual outcome.

Can we do nearly as well as best in hindsight?



We are **not** given any other info besides the yes/no bits produced by the experts. We make **no assumptions** about the quality or independence of the experts.

We cannot hope to achieve an absolute level of quality in our predictions.

Simpler question

- We have n "experts".
- One of these is perfect (never makes a mistake).
We don't know which one.
- A strategy that makes no more than $\lg(n)$ mistakes?

Halving algorithm

Take **majority vote** over all experts that have been correct so far.

I.e., if $\#$ surviving experts predicting 1 $>$ $\#$ surviving experts predicting 0, then predict 1; else predict 0.

Claim: If one of the experts is perfect, then at most $\lg(n)$ mistakes.

Proof: Each mistake cuts $\#$ surviving by factor of 2, so we make $\leq \lg(n)$ mistakes.

Note: this means ok for n to be very large.

Using "expert" advice

- If one expert is perfect, get $\leq \lg(n)$ mistakes with halving algorithm.
- But what if none is perfect? Can we do nearly as well as the best one in hindsight?

Using "expert" advice

Strategy #1: Iterated halving algorithm.

- Same as before, but once we've crossed off all the experts, restart from the beginning.
- Makes at most $\log(n) * [OPT+1]$ mistakes, where OPT is #mistakes of the best expert in hindsight.

Divide the whole history into epochs. Beginning of an epoch is when we restart Halving; end of an epoch is when we have crossed off all the available experts.

At the end of an epoch we have crossed all the experts, so every single expert must make a mistake. So, the best expert must have made a mistake. We make at most $\log n$ mistakes per epoch.

- If $OPT=0$ we get the previous guarantee.

Using "expert" advice

Strategy #1: Iterated halving algorithm.

- Same as before, but once we've crossed off all the experts, restart from the beginning.
- Makes at most $\log(n) * [OPT+1]$ mistakes, where OPT is #mistakes of the best expert in hindsight.

Wasteful. Constantly forgetting what we've "learned".

Can we do better?



Weighted Majority Algorithm

Key Point:

A mistake doesn't completely disqualify an expert.
Instead of crossing off, just lower its weight.

Weighted Majority Algorithm

- Start with all experts having weight 1.
- Predict based on weighted majority vote.
 - If $\sum_{i:x_i=1} w_i \geq \sum_{i:x_i=0} w_i$ then predict 1
else predict 0

Weighted Majority Algorithm

Key Point:

A mistake doesn't completely disqualify an expert. Instead of crossing off, just lower its weight.

Weighted Majority Algorithm

- Start with all experts having weight 1.
- Predict based on weighted majority vote.
- Penalize mistakes by cutting weight in half.

					prediction	correct
weights	1	1	1	1		
predictions	Y	Y	Y	N	Y	Y
weights	1	1	1	.5		
predictions	Y	N	N	Y	N	Y
weights	1	.5	.5	.5		

Analysis: do nearly as well as best expert in hindsight

Theorem: If M = # mistakes we've made so far and
 OPT = # mistakes best expert has made so far, then:

$$M \leq 2.4(OPT + \lg n)$$

Analysis: do nearly as well as best expert in hindsight

Theorem: If M = # mistakes we've made so far and OPT = # mistakes best expert has made so far, then:

$$M \leq 2.4(OPT + \lg n)$$

Proof:

- Analyze W = total weight (starts at n).
- After each mistake, W drops by at least 25%.
So, after M mistakes, W is at most $n(3/4)^M$.
- Weight of best expert is $(1/2)^{OPT}$. So,

$$(1/2)^{OPT} \leq n(3/4)^M$$

$$(4/3)^M \leq n2^{OPT}$$

$$M \leq 2.4(OPT + \lg n)$$

constant
ratio

Randomized Weighted Majority

$2.4(\text{OPT} + \lg n)$ not so good if the best expert makes a mistake 20% of the time.

Can we do better?



- **Yes.** Instead of taking majority vote, use weights as probabilities & predict each outcome with prob. \sim to its weight. (e.g., if 70% on up, 30% on down, then pick 70:30)

Key Point: smooth out the worst case.

Randomized Weighted Majority

$2.4(\text{OPT} + \lg n)$ not so good if the best expert makes a mistake 20% of the time.

Can we do better?



- **Yes.** Instead of taking majority vote, use weights as probabilities. (e.g., if 70% on up, 30% on down, then pick 70:30)
- Also, generalize $\frac{1}{2}$ to $1 - \epsilon$.

Equivalent to select an expert with probability proportional with its weight.

Randomized Weighted Majority

Initially: $w_1^i = 1$ and $p_1^i = 1/n$, for each expert i .

At time t :

See experts predictions. Select an expert i with probability p_t^i .

See correct answer; induces loss vector l_t .

Update the weights.

If $l_t^i = 1$, then let $w_{t+1}^i = w_t^i(1 - \epsilon)$;

else ($l_t^i = 0$) let $w_{t+1}^i = w_t^i$.

Let $p_{t+1}^i = w_{t+1}^i / W_{t+1}$, where $W_{t+1} = \sum_{i \in X} w_{t+1}^i$

Formal Guarantee for Randomized Weighted Majority

Theorem: If M = expected # mistakes we've made so far and OPT = # mistakes best expert has made so far, then:

$$M \leq (1+\varepsilon)OPT + (1/\varepsilon) \log(n)$$

Summarizing

- $E[\# \text{ mistakes}] \leq (1+\epsilon)OPT + \epsilon^{-1}\log(n)$ 🧐
- If set $\epsilon = (\log(n)/OPT)^{1/2}$ to balance the two terms out (or use guess-and-double), get bound of

- $E[\text{mistakes}] \leq OPT + 2(OPT \cdot \log n)^{1/2}$

Note: Of course we might not know OPT , so if running T time steps, since $OPT \leq T$, set ϵ to get additive loss $(2T \log n)^{1/2}$

- $E[\text{mistakes}] \leq OPT + 2(T \cdot \log n)^{1/2}$ regret

- So, $\text{regret}/T \rightarrow 0$. [no regret algorithm]

What if have n options, not n predictors?

- We're not combining n experts, we're choosing one.
- Can we still do it?
- Nice feature of RWM: can be applied when experts are n different options
 - E.g., n different ways to drive to work each day, n different ways to invest our money.

What if have n options, not n predictors?

- We're not combining n experts, we're choosing one.
- Can we still do it?
- Nice feature of RWM: can be applied when experts are n different options
 - E.g., n different ways to drive to work each day, n different ways to invest our money.
- We did not see the predictions in order to select an expert (only needed to see their losses to update our weights)

Decision Theoretic Version; Formal model

- There are n experts.
- For each round $t=1, 2, \dots, T$
 - No predictions. The learner produces a probab distr. on experts based on their past performance p_t .
 - The learner is given a loss vector l_t and incurs expected loss $l_t \cdot p_t$.
 - The learner updates the weights.

The guarantee also applies to this model!!! 

[Interesting for connections between GT and Learning.]

Can generalize to losses in $[0,1]$

- If expert i has loss l_i , do: $w_i \leftarrow w_i(1-l_i\varepsilon)$.

[before if an expert had a loss of 1, we multiplied by $(1-\text{epsilon})$, if it had loss of 0 we left it alone, now we do linearly in between]

- Same analysis as before.

Summary

- Can use to combine multiple algorithms to do nearly as well as best in hindsight.
- Can apply RWM in situations where experts are making choices that cannot be combined.
 - E.g., repeated game-playing.
 - E.g., online shortest path problem
- Extensions:
 - "bandit" problem.
 - efficient algs for some cases with many experts.
 - Sleeping experts / "specialists" setting.

Exam #2

- In-class exam on Friday, May 3rd
 - About 80% is on topics from after the midterm
 - No electronic devices
 - You are allowed to bring one $8\frac{1}{2} \times 11$ sheet of notes (front and back)

Exam #2

- All topics before the midterm (20%)
- Topics after the midterm (80%)
 - Generalization and Overfitting; Model Selection, Regularization; Error plots (training/generalization)
 - Linear Regression
 - Neural Networks. Backpropagation. Convolutional NNs
 - Boosting. Adaboost
 - Objective Based Clustering (k-means and k-means++); Hierarchical clustering (Single, complete linkage)

Exam #2

- All topics before the midterm (20%)
- Topics after the midterm (80%)
 - PCA; Dimensionality Reduction
 - Active Learning (Active SVM, Sampling bias, Safe Disagreement Based Active Learning Schemes)
 - Semi-supervised Learning (Transductive SVM; Co-training)
 - Reinforcement Learning